

Assignment2-solutions

February 5, 2022

```
[ ]: import sys
sys.path.append("../")
from dataclasses import dataclass
from typing import Mapping, Dict, Tuple
from rl.distribution import Categorical, FiniteDistribution, Constant
from rl.markov_process import FiniteMarkovProcess, FiniteMarkovRewardProcess, NonTerminal
import matplotlib.pyplot as plt
import itertools

plt.rcParams['figure.figsize'] = (15, 7)
```

1 Question 1:

The state space is: $S_t = \{1, 2, \dots, 100\}$ where 100 is the terminal state (we win the game when we reach 100).

The transition probability matrix is:

$$P_{ij} = \begin{cases} \frac{1}{6}, & \text{if state } i \text{ goes to state } j \\ 0, & \text{otherwise} \end{cases}$$

For example, if I were at state 12, then I have 6 different possibilities for the next state: 13, 14, 15, 16, 17, 18, or 19. If a ladder exists from 16 to 32, then $P_{12,32} = \frac{1}{6}$ and $P_{12,16} = 0$ instead. We consider the case of a snake in a similar manner, except that here $i > j$.

Question 2

```
[ ]: @dataclass(frozen=True)
class State:
    position : int

class SnakesAndLaddersMP(FiniteMarkovProcess[State]):
    def __init__(self, from_to : Mapping[State, State]):
        '''
        @from_to lists the ladders and the snakes.
        For example, if there was a ladder from state 43 to state 87,
        then this would be represented in the dictionary as:
```

```

        {State(position = 43) : State(position - 87)}, and so on.
        '''

        self.from_to = from_to
        super().__init__(self.get_transition_map())

    def get_transition_map(self) -> Mapping[State,
↪FiniteDistribution[State]]:
        d : Dict[State, FiniteDistribution[State]] = {}

        for state in range(1, 100):
            state_probs_map = {}
            next_state = 0

            for j in range(state + 1, min(101, state + 7)):
                if j in self.from_to.keys():
                    next_state = self.from_to[j]
                else:
                    next_state = j

            state_probs_map[State(position = next_state)] = ↪
↪1. / 6.

            # If I'm at state 97, I need only a 3 to reach 100.
            # If I get more than a 3, I stay in the same position.
            if state > 94:
                state_probs_map[State(position = state)] = ↪
↪(state - 94.) / 6.

            d[State(position = state)] = ↪
↪Categorical(state_probs_map)

        return d

```

1.0.1 Create an instance of the Snakes and Ladders game

```
[ ]: changes_from = [1, 4, 9, 28, 36, 21, 51, 71, 80, 16, 47, 49, 56, 64, 87, 93, ↵
    ↪95, 98]
changes_to = [38, 14, 31, 84, 44, 42, 67, 91, 100, 6, 26, 11, 53, 60, 24, 73, ↵
    ↪75, 78]

from_to = {fr : to for fr, to in zip(changes_from, changes_to)}
game = SnakesAndLaddersMP(from_to = from_to)

# print("Transition Map")
# print("-----")
# print(game)
```

1.0.2 Generate the constant distribution of the starting states and generate traces of the game and count the total number of dice rolls needed per game to finish it

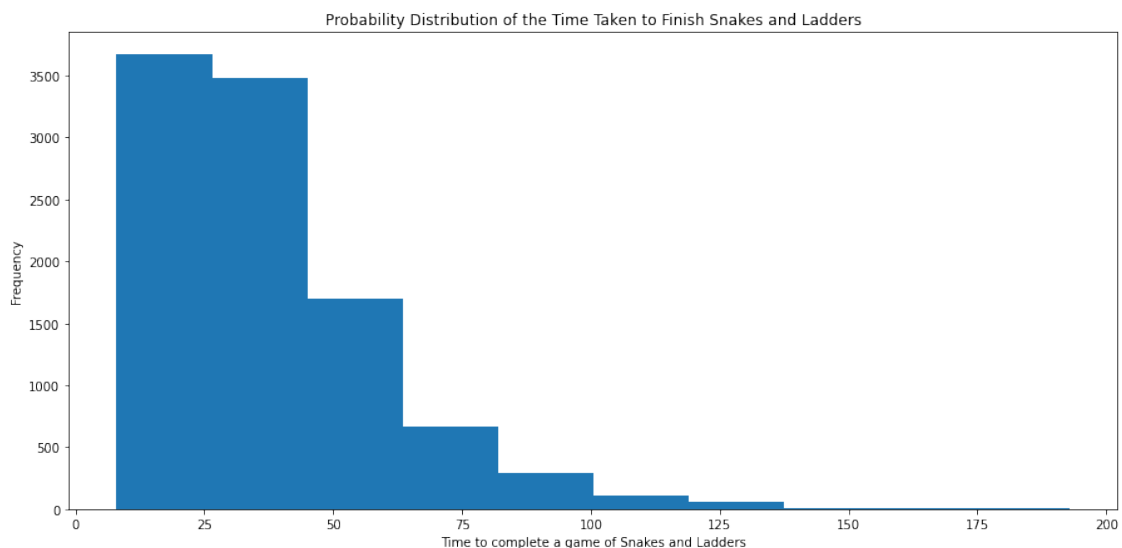
```
[ ]: start_distribution = Constant(value = NonTerminal(State(position = 1)))
num_traces = 10000

outcomes = [len([st for st in it]) for it in itertools.islice(game.
    ↪traces(start_distribution), num_traces)]
```

1.0.3 Plot the histogram of the time steps (number of dice rolls)

```
[ ]: plt.hist(outcomes)

plt.xlabel('Time to complete a game of Snakes and Ladders')
plt.ylabel('Frequency')
plt.title("Probability Distribution of the Time Taken to Finish Snakes and ↵
    ↪Ladders")
plt.show()
```



2 Question 3

```
[ ]: @dataclass(frozen=True)
class State:
    position : int

class FrogPuzzle(FiniteMarkovProcess[State]):
    length : int = 0 # this is the length of the river between the two banks

    def __init__(self, length : int):
        self.length = length
        super().__init__(self.get_transition_map())

    def get_transition_map(self) -> Mapping[State,
↪FiniteDistribution[State]]:
        d : Dict[State, FiniteDistribution[State]] = {}

        for state in range(1, self.length + 2):
            state_prob_map = {}
            for next_state in range(state + 1, self.length + 3):
                state_prob_map[State(position=next_state)] = 1.
↪/ (self.length - state + 2)

            d[State(position = state)] = Categorical(state_prob_map)

        return d
```

2.0.1 Create an instance of the Markov process

```
[ ]: L = 20

puzzle = FrogPuzzle(L)
# print("Transition Map")
# print("-----")
# print(puzzle)
```

2.0.2 Generate traces of it and get the total number of jumps required to get to the other side of the river

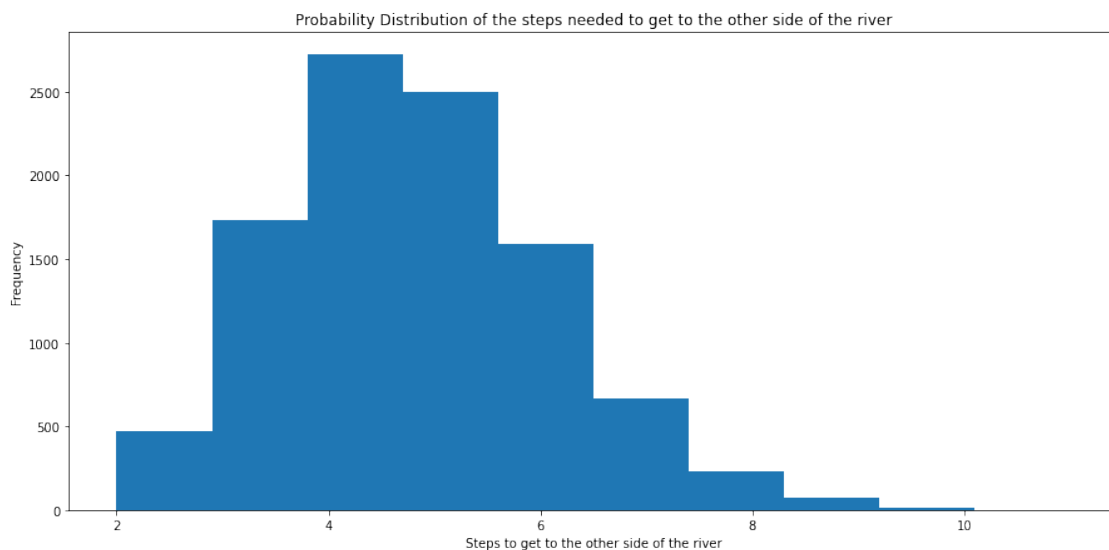
```
[ ]: start_distribution = Constant(value = NonTerminal(State(position = 1)))
num_traces = 10000
```

```
outcomes = [len([i for i in trace]) for trace in itertools.islice(puzzle.
    ↳traces(start_distribution), num_traces)]
```

2.0.3 Plot a histogram of the jumps required to get to the other side of the river

```
[ ]: plt.hist(outcomes)

plt.xlabel('Steps to get to the other side of the river')
plt.ylabel('Frequency')
plt.title("Probability Distribution of the steps needed to get to the other_
    ↳side of the river")
plt.show()
```



3 Question 4

Here I recreate the Snakes and Ladders game, but with rewards associated with each state. In order to find the expected number of dice rolls to finish a game, I set the rewards of all states to 1.

```
[ ]: @dataclass(frozen=True)
class State:
    position : int

class SnakesAndLaddersMRP(FiniteMarkovRewardProcess[State]):
    def __init__(self, from_to : Mapping[State, State]):
        self.from_to = from_to
        super().__init__(self.get_transition_reward_map())
```

```

def get_transition_reward_map(self) -> Mapping[State,
↪FiniteDistribution[Tuple[State, float]]]:
    d : Dict[State, FiniteDistribution[Tuple[State, float]]] = {}
    reward = 1

    for state in range(1, 100):
        state_probs_map = {}
        next_state = 0

        for j in range(state + 1, min(101, state + 7)):
            if j in self.from_to.keys():
                next_state = self.from_to[j]
            else:
                next_state = j

        state_probs_map[(State(position = next_state),
↪reward)] = 1. / 6.

        if state > 94:
            state_probs_map[(State(position = state),
↪reward)] = (state - 94.) / 6.

        d[State(position = state)] =
↪Categorical(state_probs_map)

    return d

```

3.0.1 Create an instance of the game

```

[ ]: changes_from = [1, 4, 9, 28, 36, 21, 51, 71, 80, \
                    16, 47, 49, 56, 64, 87, 93, 95, 98]
changes_to = [38, 14, 31, 84, 44, 42, 67, 91, 100, \
             6, 26, 11, 53, 60, 24, 73, 75, 78]

from_to = {fr : to for fr, to in zip(changes_from, changes_to)}
game = SnakesAndLaddersMRP(from_to = from_to)
gamma = 1.0 # the discount factor

# print("Transition Map")
# print("-----")
# print(FiniteMarkovProcess({s.state: Categorical({s1.state: p for s1, p in v.
↪table().items()}) \
#                                     for s, v in game.transition_map.items()}))

# print()
# print("Transition Reward Map")
# print("-----")

```

```

# print(game)

# print()
# print("Reward Function")
# print("-----")
# game.display_reward_function()

# print()
# print("Value Function")
# print("-----")
# game.display_value_function(gamma = gamma)

```

3.0.2 Use Monte Carlo simulations to find the expected number of dice rolls to win the game and compare this with the value obtained using Markov reward processes:

```

[ ]: val = game.get_value_function_vec(gamma = gamma)
print(f"The expected number of dice rolls (using the value function) is {val[0]:
    ↪.3f}.".")

start_distribution = Constant(value = NonTerminal(State(position = 1)))
num_traces = 10000
outcomes = [len([st for st in it]) for it in itertools.islice(game.
    ↪reward_traces(start_distribution), num_traces)]
print(f"The expected number of dice rolls (using Monte Carlo)          is
    ↪{sum(outcomes) / num_traces:.3f}.".")

```

The expected number of dice rolls (using the value function) is 36.819.

The expected number of dice rolls (using Monte Carlo) is 36.787.